


Pure LLM Naive RAG LightRAG Full system
L1 literature-consensus benchmark (118 cases): baselines under two output budgets

(a) 10-token, single call



(b) 3000-token, reasoning model, majority of 5

